

UESTC_00 ASC 2018 Proposal

Abstract

This paper presents the corresponding scheme and test results for the four topics of the ASC-2018 contest.

First of all, this paper presents the efforts and achievements made by the University of Electronic Science and Technology in supercomputers, and the basic conditions of the participating teams. Then, we presented the HPC platform we designed under the 3000 watt power consumption limit. And we have already Hardware-based HPL and HPCG-based performance tuning. Next, this paper presents optimized parameters and code adjustments for Relion and gives the heaviest results. Finally, we give the neural network we designed for Answer Prediction for Search Query's Models and results.

Contents

1	Supercomputing Activities in UESTC	3
1.1	Background	3
1.2	Supercomputing-Related Hardware and Software Platforms	3
1.2.1	Hardware Platforms	3
1.2.2	Software Platforms	4
1.3	Supercomputing-Related Courses, Trainings, and Interest Groups	4
1.4	Supercomputing-Related Research and Applications	5
1.4.1	High-level Laboratory planning and construction	5
1.4.2	Supercomputing Laboratory Service College Student Competition	5
1.4.3	Teaching and Scientific Research Promote Each Other	5
1.5	The Key Achievements on Supercomputing Education	5
1.5.1	Our HPC Education Development Process	5
1.5.2	The Status of HPC Education in Our School	6
2	Introduction of Our Team	6
2.1	Overview	6
2.2	Team Members	7
2.3	Group Photos	12
2.4	Team Slogan	12
3	HPC System Design	12
3.1	Hardware Environment	12
3.2	Software Environment	14
4	HPL Test	15
4.1	Hardware Platform and Theoretical Performance	15
4.2	Software Platform	15
4.3	Optimization Process	15
4.3.1	Round 0	15
4.3.2	Round 1	16
4.3.3	Round 2	17
4.3.4	Round 3	18

4.3.5	Round 4	18
4.4	Summary	18
5	HPCG Test	20
5.1	Hardware Platform	20
5.2	Software Platform	20
5.3	Summary	20
6	The RELION Test	21
6.1	Hardware Platform	21
6.2	Software Platform	21
6.3	Parameter Tuning	21
6.4	Code Optimization	22
6.4.1	First Hot Point	22
6.4.2	Second Hot Point	22
6.4.3	Third Hot Point	23
6.5	Answer Summary	23
6.5.1	Step I : Class 2D	23
6.5.2	Step II : Class 3D	24
6.5.3	Step III : Refine 3D	24
7	Answer Prediction for Search Query	24
7.1	Hardware Platform	24
7.2	Software Platform	25
7.3	Summary	25
	Appendices	26

1 Supercomputing Activities in UESTC

1.1 Background

Due to chip power consumption and physical limitations, Moore's Law continues to be effective and results in higher processor performance. The traditional method of increasing the processor frequency has shifted to increasing the number of computing cores. The processor rapidly evolves from single core to multi-core and multiple cores. In fact, mainstream personal processor chips have virtually halted the production of single-core processors since 2007. Commercial CPUs currently on the market usually integrate 4 to 8 computing cores, complemented by the hyper-threading technology provided by the chip manufacturers (such as Intel HT), enabling parallel computation of 8 to 16 physical threads. The introduction of GPGPU products, making the PC's computing power to further strengthen. Computer system architecture has been fully entered the era of multi-core processors from the era of single-core processors.

In the multi-core, all-nuclear chip-led consumer market, parallel computing is no longer super-computing proprietary technology. Parallel computing has expanded from the field of supercomputing to traditional business applications and personal information services and has become the core technology to improve the energy efficiency and energy saving of computer programs.

Supercomputing is the frontier research field of current computer science and technology, which is a high-level application of parallel computing. It is of strategic importance to the national economy, national security and the progress of science and technology, and is one of the important symbols to measure the overall strength of a country. In the recent ten years, the global supercomputer jumped from ten trillion times to several petaflops, achieving a hundredfold increase in performance. It is estimated that a terabyte-scale supercomputing platform will be built around 2018, which will be 20 times faster than the fastest Tianhe II now in the world to meet the needs of many fields including astrophysics and cosmology, physics and chemistry, life sciences, earth sciences, Systems and Engineering Science in five major areas of science and engineering computing needs.

1.2 Supercomputing-Related Hardware and Software Platforms

1.2.1 Hardware Platforms

Inspur Tiansuo series HPC platform (Show in table 1).

Table 1: Hardware Platforms in UESTC

Device Name	Number
Inspur TS4220 Cabinet	1
Inspur S1000KVM Console	1
NF5280M4 Management Node	1
NF5280M4 GPGPU Compute Node	8
H3C S1224 Gigabit Ethernet Switch	2
Intel Xeon Phi 7120P Coprocessor	6
Nvidia Tesla K80	4
GForce GTX 780	6

1.2.2 Software Platforms

- Operating System: RedHat Linux
- Development Kit
 - Intel C++/Fortran Compiler
 - Intel MKL Math Core Library
 - GNU series compiler
- Operating environment
 - MVAPICH2, OpenMPI and MPICH Parallel Environment
 - HPC, Hadoop and Spark Computing Environment
- Tuning Documentation
 - MVAPICH2-X/GDR/MIC, HPC-X Toolkit, OpenMPI and MPICH2 Parameter Tuning Documentation
 - ATLAS, ApenBLAS, Lapack, ScalAPACK and FFTW3 Mathematical Library Parameter Tuning Documentation

1.3 Supercomputing-Related Courses, Trainings, and Interest Groups

- Object-Oriented Programming Java: JavaThreads multi-threaded programming.
- Operating system, computer network programming: PThreads multi-threaded program development.
- Distributed and Parallel Computing, Numerical Analysis, Computational Math Fundamentals: MPI, OpenMP, CUDA Programming, OpenCL Heterogeneous Programming.
- CUDA programming: CUDA.

1.4 Supercomputing-Related Research and Applications

The supercomputing project leader Guoming Lu has more than 10 years of experience in supercomputing and cloud computing. He has led students to participate in college students' supercomputing competitions for many times, and has a deep background in scientific research and teaching in areas related to supercomputing. The other members of the team are also related to distributed computing and data mining. At the same time, the laboratory has completed several research projects and experimental projects in supercomputing and high performance computing. Therefore, the team has a good research in supercomputing. And teaching basis.

1.4.1 High-level Laboratory planning and construction

Project Leaders successfully applied for NSF "Parallel Distributed Computing Curriculum Innovation Project" and NVIDIA "CUDA Teaching Center" in October 2013 and October 2014, respectively, through experts in the field of supercomputing and world-class GPUs. Hardware manufacturers communicate and communicate to ensure a high level of laboratory planning and experimental design.

1.4.2 Supercomputing Laboratory Service College Student Competition

As a supercomputing laboratory, the supernumerary competition for students has been trained and occupied. The laboratory is committed to greatly improving the college students' supercomputing level, which will enable UESTC to occupy a place in the international arena of the supercomputing competition.

1.4.3 Teaching and Scientific Research Promote Each Other

Through the supercomputing lab, a school's own supercomputing experimental computing platform, service computer college, and other colleges' supercomputing requirements were established. During the process of in-depth scientific research, extensive application cases were collected, and the supercomputing experiment was extensively explored and enriched. Achieve mutual promotion of scientific research and teaching.

1.5 The Key Achievements on Supercomputing Education

1.5.1 Our HPC Education Development Process

Our school of computer science in the response to the trend of computer parallel professional training program to adjust in most domestic colleges and universities in front of. In 2010, the School of Computer Science and Technology of the UESTC revised the program of computer science (undergraduate) training and opened the "distributed parallel computing" course (the seventh semester) at the 2010 training program to adapt to the traditional training of multicore and parallel computing technologies Program

challenges. In October 2013, as the third university in mainland China, UESTC successfully joined the American NSF "Parallel Distributed Computing Curriculum Innovation Project", initiated the curriculum reform of computer science and technology, and began to strengthen the overall reform of professional courses. Parallel computing is taught in various courses in computer science and technology. The project initially covered four topics: Computer Architecture, Computer Network Programming, Java Programming Language Design and New Open Course Distributed and Parallel Computing, focusing on the research of distributed parallel computing knowledge points in the course group Reasonable distribution, to help students from hardware to software to establish the concept of distributed parallel computing, to help students adapt to the development trend of computer parallelism. At present, distributed parallel computing courses were first taught to seniors in the fall of 2013 and were well received by students. According to the results of the pre-course reform, the School of Computer Science put forward that the three courses of algorithm and data structure, operating system and compilation principle are also included in the teaching reform curriculum group so as to provide students with comprehensive knowledge of distributed parallel computing and establish complete knowledge of parallel computing system. In October 2014, the course reform program was once again nominated for NSF's "Parallel Distributed Computing Curriculum Innovation Project", and the University of Electronic Science and Technology became the only university to enter the program twice (Fall 2013, Fall 2014). Our school was short-listed in the curriculum reform project twice, providing professional support and guidance for the curriculum reform in computer science in our university.

1.5.2 The Status of HPC Education in Our School

In April 2014, the project leader was successfully supported by NVIDIA-sponsored CUDA teaching center project and became the second CUDA teaching center (CTC) in southwest China to conduct GPGPU-based supercomputing and heterogeneous computing teaching for UESTC. And practice provides a first-rate platform. NVIDIA is the world's largest supplier of GPUs, the CUDA Teaching Center launched in 2011, which supports colleges and universities in providing CUDA teaching materials, software licenses, GPGPU donations and technical services. This project not only provides support for the construction of first-rate teaching environment in universities and colleges, but also provides practical exercises for students during CUDA learning so as to achieve the goal of teaching-learning. From 2011 Plan Implementation So far, internationally renowned universities including MIT, Carnegie Mellon University, University of California Fallon, Purdue University, Michigan State University and University of Toronto have joined the project as CUDA-based supercomputer Teaching provides a wealth of teaching support.

2 Introduction of Our Team

2.1 Overview

Through the education program of the School of Computer Science and Engineering at the University of Electronic Science and Technology, the University of Electronic Science and Technology has established a series of courses and competition training for

undergraduates. Supercomputers are one of the important directions for the development of computers in the future, and are also important tools for the development of many fields. Based on this, we formed a team to participate in the ASC Student Supercomputer Challenge competition in 2018. (Our team is one of the two teams participating in the University of Electronic Science and Technology in 2018).

2.2 Team Members

Team advisor : Wenhan Zhan

University of Electronic Science and Technology of China(UESTC)

Education Background

- 2014~, Ph.D Candidates of Software Engineering, UESTC
- 2010~2013, Master of Computer System Architecture, UESTC
- 2006~2010, Bachelor of Software Engineering, UESTC

Work Experience

- 2015~, Experimentalist of National Demonstration Center for Experimental Computer Education, UESTC
- 2013~, Director of Student Innovation Center of Computer Science and Engineering, UESTC
- 2013~2015, Teaching assistant of School of Computer Science and Engineering, UESTC

Scholarship and Award

- 2015, Third Prize of Sichuan Science and Technology Progress Award, Sichuan Provincial Government
- 2015, Third Prize of University Quality Education Outstanding Brand Activities, China Higher Education Society
- 2016, First Prize of Teaching Achievement Award, UESTC
- 2006~2013, Outstanding Graduate of UESTC, Outstanding Graduate Student of Sichuan Province, National student scholarship, "Google" student scholarship, "Samsung" student scholarship, etc.

Research Project

- Development of Audio and Video Big Data Identity Recognition System and Livelihood Applications. 2,000,000

- Industrial Cloud Image Recognition and Production Line Forecasting . 1,750,000
- Distribute Business Integrated Platform . 1,000,000

Research Paper

- A high-performance distributed file system for large-scale concurrent HD video streams, *Concurr. Comput.* , vol. 27, no. 13, pp. 3510-3522, 2015.
- CSTORE: A desktop-oriented distributed public cloud storage system, *Comput. Electr. Eng.*, vol. 42, pp. 60-73, 2015.
- CSTORE: A massive user-oriented file system with high scalability, *J. Comput. Inf. Syst.*, vol. 10, no. 16, 2014.
- Three-level mapping hash method in large-scale distributed storage system, *Int. J. Adv. Comput. Technol.*, vol. 4, no. 19, 2012.
- A scalable data distributing method supporting replication and weight, in *Proceedings - 2012 8th International Conference on Computing Technology and Information Management, ICCM 2012*, 2012, vol. 1.

Authorized Patent

- Large scale distributed storage system with replication based on multi-mapping and its application method. No: ZL201210083042.7.
- Method for implementing distributed storage system based on load adjustment and fault-tolerant. No: ZL201210172660.9.

Captain : Shengsong Tan

School of Computer Science and Engineering, University of Electronic Science and Technology, Major in Digital Media. Specializing in Algorithm Design and Deep Learning. Some awards are as follows:

- Silver Medal, The 2017 ACM-ICPC Asia Urumqi Regional Contest Dec. 2017
- Silver Medal, The 2017 ACM-ICPC Asia Beijing Regional Contest Nov. 2017
- First Prize, The 10th National College Student Information Security Contest Jul. 2017
- Bronze Medal, The 30th National Olympiad in Informatics (NOI) Jul. 2013
- Bronze Medal, The 29th National Olympiad in Informatics (NOI) Jul. 2012

Experience

- Swarmart Co.Ltd Start-up company

- Focused on crowdsourcing annotation
 - Developed algorithms on crowdsourcing platform
 - Built a quality control system on crowdsourcing platform to control tasks distribution and result computation
- A Multi-Factor Authentication System
 - It mixed facial features, voiceprint, and lips action to achieve high accuracy in identification and verification
 - Led the team to achieve a high accuracy of the system and won the first prize in final contest
 - Employed deeplearning method to extract facial features
 - Detected face liveness by tracing lips action
 - Designed the GUI using Python-Qt
 - Supervised the team to implemented a voiceprint module and wrote project documents
 - link: <https://github.com/tsstss123/faceUnionVoiceRecognition>
- Chinese CAPTCHA Recognition Tool
 - Designed and trained a suitable convolutional neural network in MXNet targeting two specific types of Chinese CAPTCHA
 - Dealt with other work including image preprocessing and training dataset collection
 - Achieved an accuracy rate over 85% which meets the standard of our customer
- A Simple Ray Tracing Demo May
 - Ranked first among all the individual projects of the advanced computer graphics course
 - Implemented collision detection methods based on bounding volume hierarchies to accelerate rendering
 - link: <https://github.com/tsstss123/RayTracing>

Member : Zihui Liang

Computer Science and Engineering, University of Electronic Science and Technology, Information Security, a freshman student. Specializes in algorithm analysis and design. Some awards are as follows:

- NOIP2016 first prize
- GDKOI2016 First Prize (City Election)
- GDOI2016 Second Prize (Provincial Election)
- APIO (Asia Pacific Informatics China Division) 2016 Silver

- NOI2016 silver medal
- NOI Winter Camp 2016 Gold Medal
- Tsinghua Summer Camp 2nd agreements
- Twelfth CCF CSP Certification Exam National Top 9
- 2017 Baidu Star Online top 500
- 2016 Google Code Jam Global Top 3000
- IEEE Extreme Programming Contest Top 10
- Codeforces Round #429 (Div 2) rank 1
- Codeforces candidate master

Focus on algorithm learning, design, and network security related orientation learning.

Member : Zitong Li

University of Electronic Science and Technology, Computer Science and Engineering College of Computer Science and Technology is studying sophomore undergraduate students. Some awards are as follows:

- First prize of National Informatics Olympic League (NOIP2014)
- National Informatics Olympiad (NOI2015) Bronze Award
- 2016 ACM-ICPC Asian Regional Competition (Shenyang Station) Silver Award
- 2017 AUC-ICPC Asian Regional Competition (Qingdao Station) Bronze Award
- 2017 AUC-ICPC Asian Regional Competition (Nanning Station) Silver Award
- 2017 China College Programming Contest (CCPC) (China Collegiate Programming Contest) (Harbin Station) Silver Award
- The Gold Medal of Sichuan College Student Program Design Competition 2017
- The First Prize of the China Collegiate Computing Contest - Group Programming Ladder Tournament Team in 2017 China University Computer Competition (CCCC)
 - Group Programming Ladder Tournament (China Collegiate Computing Contest
 - Group Programming Ladder Tournament)

There have also been learning experiences in mathematical modeling related courses. The research on algorithms, data structures, and other aspects is more in-depth, and it is used more often for programming languages such as C, C++, and JAVA, and the object-oriented programming ideas are understood, and a certain scale of project code can be realized. There are learning experiences of MPI, openmp and other application programming interfaces that are widely used in parallel computing, and can implement small-scale parallel programs.

Member : Xiongyu Zhu

Major in Information Security, School of Computer Science and Engineering, UESTC. Interested in bleeding-edge security technologies. Part of the winning as follows :

- Third prize of IEEEExtreme Campus Qualifying Competition
- Third prize of ACM Campus & Southwestern Colleges Competition

Experience

- Swarmart Co.Ltd Start-up Company
 - DevOps Engineer
 - Built the Microservice Architecture
 - Developed a Backend Service
- Multiple Commercial Projects
 - Deployment of various services
 - Load Balance
 - Continuous Integration
 - Years of experience in different Linux distributions

Member : Yanjing Ren

Graduate of Information Security at the School of Computer Science and Engineering at the University of Electronic Science and Technology of China. Experienced in mathematical modeling and large-scale storage system development. Has a certain understanding of data structure and algorithm analysis. Learn about C++, Python, Matlab object-oriented programming. Hadoop studied systematically. With the Spark framework, it is able to achieve large-scale parallel computing program design, and is better at multi-threaded and multi-process programming. Some of the awards are as follows:

- 2017 American Undergraduate Mathematical Modeling Contest Meritorious Winner
- 2017 National University Student Mathematical Modeling Competition

Experience

- Bucket Attack Based on FSL DataBase
 - link: <https://github.com/tinoryj/BucketAttack>
- Simple Web Server Based on Epoll & C++
 - link: https://github.com/tinoryj/SimpleHTTPServer_epoll_threadPoll

- Convergent Dispersal Deduplication Datastore
 - link: <https://github.com/tinoryj/CDStore>
- Rekeying-aware Encrypted Deduplication Storage
 - link: <https://github.com/tinoryj/REED-system>

2.3 Group Photos




Figure 1: Our Team Members

2.4 Team Slogan

If you think you can, you can.

3 HPC System Design

Objective: Design a system to achieve the best computing performance under 3,000-watt power consumption. Specify your system's software and hardware configuration and interconnection. Describe the power consumption, evaluate the performance, and analyze the advantages and disadvantages of your proposed architecture.

3.1 Hardware Environment

Our proposed system is based on the Inspur NF5280M5 server following instructions and recommendations. The components and power consumption estimation listed in the table below.

Table 2: Hardware Configuration In Our Design

Item	Name	Configuration	Power Estimation
Server	Inspur NF5280M5	CPU : Intel Xeon Platinum 8176M $\times 2$ (2.1GHz \sim 3.8GHz, 28 Cores) TDP : 165 W $\times 2$	One Node : 1403.2W (2 Node in System)
		GPU : NVIDIA TESLA V100 $\times 4$ TDP : 250 W $\times 4$	
		Memory: 64GB , DDR4 , 2400MHz $\times 12$ TDP : 6 W $\times 12$	
		Hard Disk: Intel SSD Pro 7600p $\times 24$ (512GB, M.2, PCIe 3.0x4) TDP : 50 mW $\times 24$	
HCA Card	FDR	InfiniBand Mellanox ConnectX-3 HCA card, Single port, QSFP, FDR IB	One Node : 9 W (2 Node in System)
Switch	FDR-IB Switch	<i>SwitchXTM</i> FDR InfiniBand switch, 36 QSFP port	130 W
Cable	InfiniBand cable	InfiniBand FDR optical fiber cable, QSFP port, cooperating with the InfiniBand switch for use	NaN

Within the power limitation of 3000 Watts, the configuration can be used.

Through practical experience, it is found that accommodating more computing resources as much as possible in a single node is more conducive to improving the performance/power ratio of the entire HPC platform. At the same time, the huge I/O overhead due to network transmission can be minimized to achieve Maximizing performance under 3000-watt power limitations. In addition, CUDA-based GPGPUs have far greater parallel computing power than CPUs with the same power consumption. Based on this, our design criteria is to make full use of the resources of a single node as much as possible. In the whole system, we use two NF5280M5 servers based on the configuration shown in the table 2. The theoretical power consumption of the entire platform is 2954.4 KW. When calculating the CUGPU-based GPGPU computing performance only, theoretical single-precision floating-point performance can reach 56 Tflops.

In each node, we make full use of the maximum memory capacity supported by the Intel Xeon Platinum 8176M Processors and use a single 64GB DDR4-ECC memory to increase the memory capacity as much as possible so as to use the high-speed access performance of memory to reduce system I/O bottlenecks. We use 4 blocks. On the premise of the same power consumption, the NVIDIA TESLA V100 GPU enhances the performance of the system in scenarios such as deep learning. At the same time, we use 24 NVME-based SSDs to form a RAID array to provide more than 10 GB/s of disk I/O. Throughput.

In the connection part of the two nodes in the system, we connect through 16 linked InfiniBand methods to achieve a data exchange speed of up to 8 GB/s between the two computing nodes. At the same time, the master node is accessed through the ordinary Ethernet connection for the environment configure and monitoring. The overall struc-

ture of the system is shown in Figure 2

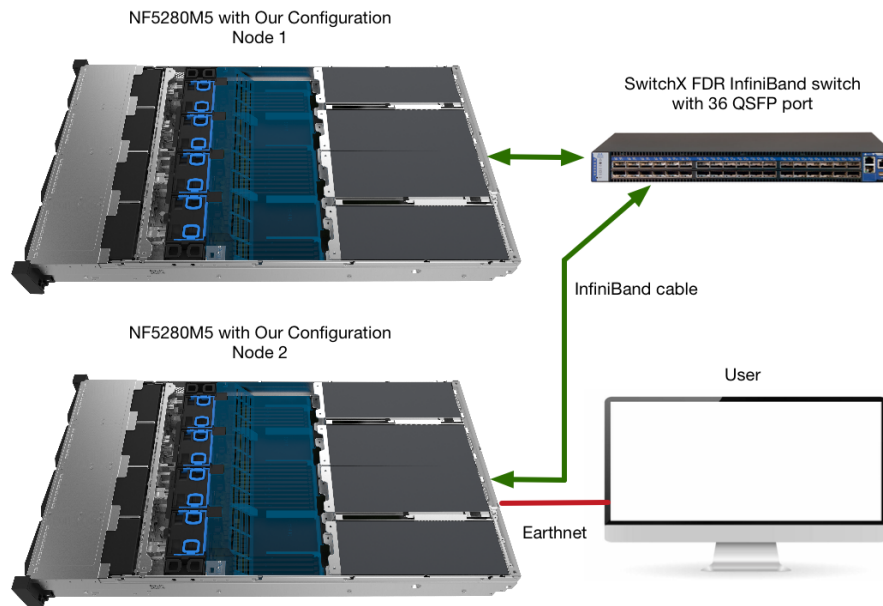


Figure 2: System Overall Structure

3.2 Software Environment

- Operating System: RedHat Linux
- Development Kit
 - Intel C++/Fortran Compiler
 - Intel MKL Math Core Library
 - GNU series compiler
 - GNU gprof
 - Intel VTUNE
- Operating environment
 - MVAPICH2, OpenMPI and MPICH Parallel Environment
 - HPC, Hadoop and Spark Computing Environment
 - NFS and RamDisk
- Cluster Monitoring and Scheduling System
 - PBS
 - OOOZIE

4 HPL Test

4.1 Hardware Platform and Theoretical Performance

We use the single-node dual-circuit computing platform for HPL tuning in this question, for our hardware platform:

- CPU : Intel Xeon E5460 (2.66GHz, 8C16T) $\times 2$
- Memory : 8GB DDR3-ECC 1600MHz $\times 18$

According to Intel official data^[4], we calculated the theoretical performance of our hardware platform to be $2.66 \times 16 \times 2 = 85.12Gflops$

4.2 Software Platform

For baseline, we use the intel HPL-2.1 in round 0. Then we use the recompiling version HPL-2.2 for optimization. The software platform is shown flow:

- **System** : Ubuntu 16.04.3 LTS
- **Compiler & MKL** : Intel Parallel Studio XE 2018.1
- **Round 0** : Intel original HPL-2.1
- **Round 1-4** : Recompiling version of HPL-2.2

4.3 Optimization Process

We first tested the calculation speed of baseline in round 0. According to Intel and HPL related documents^[1,2,3] the possible optimal parameters were calculated at $N = 130000$, $NB = 120$, $P = 4$, $Q = 4$. Afterwards, we find the optimal parameters by adjusting the four different parameters.

According to the consulted documents, we conclude that the parameters that have obvious influence are N, NB, P, and Q. According to the formula 1, the theoretical maximum performance value is $N = 135092$. The NB value ranges from 32 \sim 256. .

$$N \times N \times 8 = 80\% \times MaxMemorySize \quad (1)$$

4.3.1 Round 0

In round 0, we use the default Intel HPL environment for testing. The main parameters are shown in the following table 3:

In this round, the computing power of our computing platform that we tested floats around 38 Gflops. Some of the test results are shown in the following table 4. The final log files are detailed in the appendix.

Table 3: Parameters in Round 0 for Baseline

Parameters	Content		
Ns	130000	102400	150000
NBs	110	120	130
Ps	4		
Qs	4		

Table 4: Part of Log in Round 0

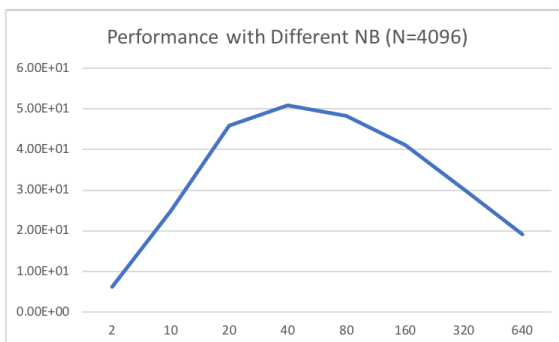

Column	Fraction	Kernel	Mflops
022110	0.170	37958.81	38495.56
022770	0.175	38081.28	38479.87
023430	0.180	38564.32	38480.45
024090	0.185	37967.46	38464.89
024750	0.190	38317.15	38461.91
025410	0.195	38406.03	38459.64
026070	0.200	37941.81	38447.85

4.3.2 Round 1

In the current round, we adjusted the values of N , NB , P , and Q within a certain range, and conducted tests to determine the actual performance under different parameters.

Influence of different NB values

We tested the computational performance of $NB = 2, 10, 20, 40, 80, 160, 320$, and 640 for the two cases $N = 4096$ and $N = 20480$. As shown in figure 3, the performance curve for different NB values when $N = 4096$. As shown in figure 4, the performance curve for different NB values when $N = 20480$. Here $P = 2$, $Q = 8$. It can be seen that as the NB value becomes larger, the performance increases first and then decreases, and there is a peak. When $NB = 80$ and $NB = 160$, the calculation performance is equivalent.

Figure 3: Performance Changes with $N = 4096$ Figure 4: Performance Changes with $N = 20480$

Influence of different N values

We tested the computational performance of $N = 2048, 4096, 10240, 20480, 40960, 102400$ in two cases $NB = 20$ and $NB = 160$. As shown in figure 5, the performance curve of different N values at $NB = 20$. As shown in figure 6, the performance curve for different N values when $NB = 160$. Here $P = 2, Q = 8$. It can be seen that the performance increases gradually as the value of N increases.

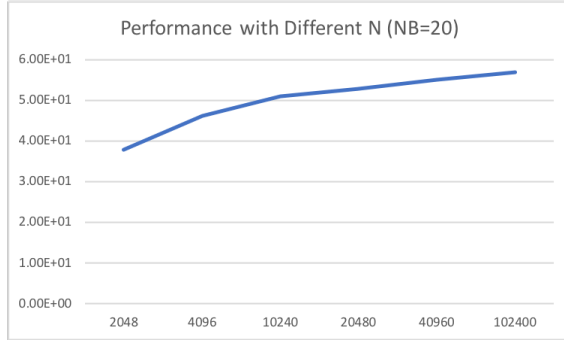


Figure 5: Performance Changes with $NB = 20$




Figure 6: Performance Changes with $NB = 160$

Influence of different $P \times Q$ values

We have tested $NB = 20, NB = 160$ and $N = 4096$, and $N = 20480$. $P \times Q = 1 \times 16, 2 \times 8, 4 \times 4, 2 \times 4, 8 \times 2$, Computational performance of these five cases. As shown in figure 7, for $NB = 20$ different N and $P \times Q$ value performance changes Curve. As shown in figure 8, the performance curve of different N and $P \times Q$ values for $NB = 160$.




Figure 7: Performance Changes with $P \times Q$ ($N = 4096$)




Figure 8: Performance Changes with $P \times Q$ ($N = 20480$)

The orange line represents $NB = 20$ and the blue line represents $NB = 160$. Through analysis, it can be seen that $P \times Q$ can optimize performance at 4×4 .

4.3.3 Round 2

Through Round 1's analysis, we found that the parameter that has the greatest impact on performance is NB . When $P \times Q = 4 \times 4$, the best performance occurred. Therefore, in this round, we selected $N = 40960, P \times Q = 4 \times 4$, and tested the calculations for $NB = 80, 160, 320$, and 640 . Performance (as shown in figure 9).

4.3.4 Round 3

Based on the previous round of testing, we determined the range of NB's value to $80 \leq NB \leq 240$, and tested the calculation performance for $NB = 80, 120, 160, 240$ (as shown in figure 10).





Figure 9: Performance Changes in Round 2 Figure 10: Performance Changes in Round 3

4.3.5 Round 4

Based on the previous round of testing, we set the NB's value range to $100 \leq NB \leq 140$. We tested the performance of calculations for $NB = 100, 110, 120, 130, 140$ (as shown in figure 11).


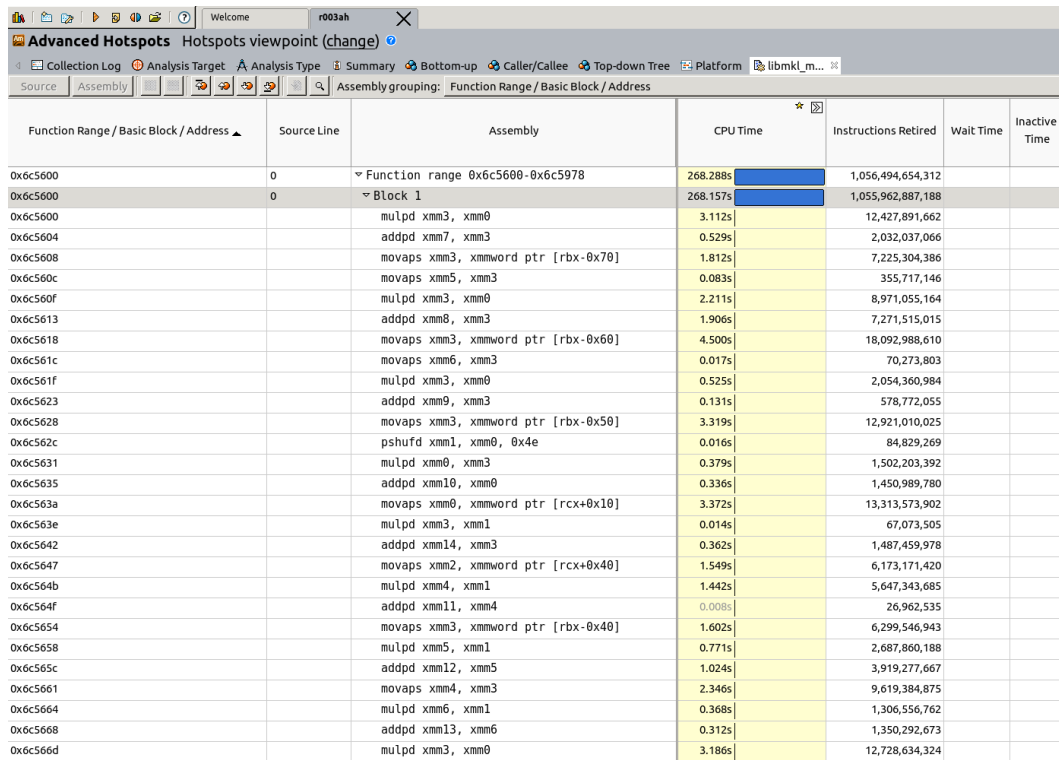


Figure 11: Performance Changes in Round 4

4.4 Summary

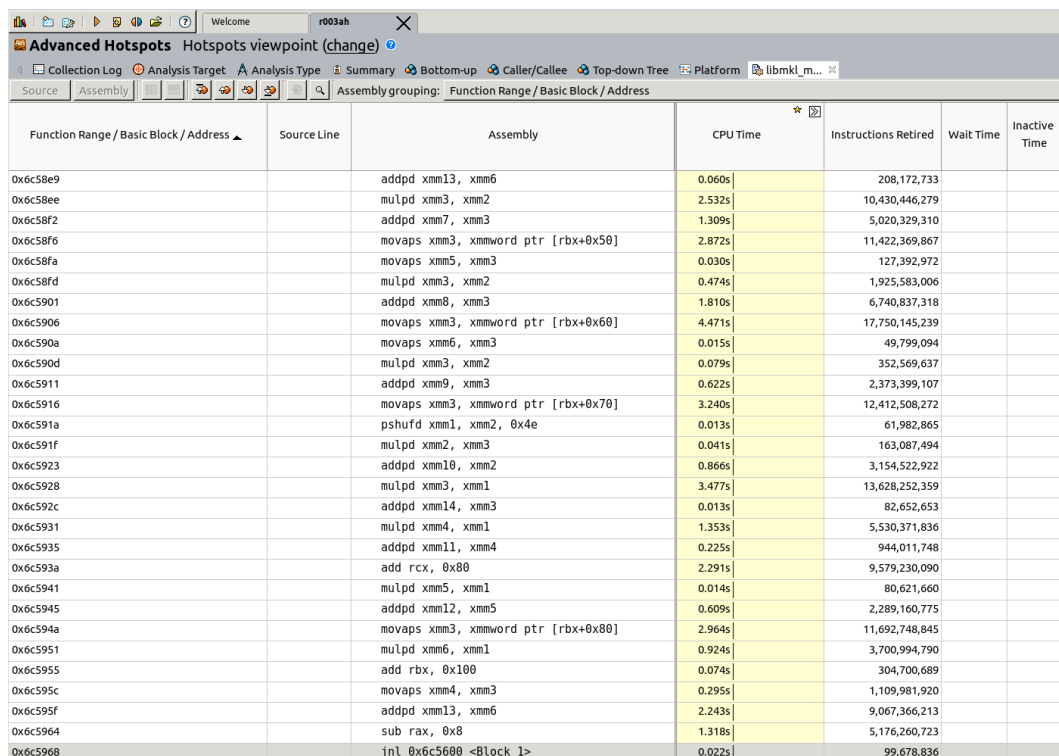
In the HPL optimization, we consulted the document and parameter optimization. By analys the profiling result against xhpl, we found that 75% time of the time was spent on one function of intel MKL, to be more specific, it's the code segment between 0x6c5600 and 0x6c5968 of libmkl_mc3.so(show in figure 12, 13).



The screenshot shows the 'Advanced Hotspots' tool interface with the 'Hotspots viewpoint (change)' selected. The 'Assembly' tab is active, displaying a table of assembly instructions grouped by function range. The table includes columns for Function Range / Basic Block / Address, Source Line, Assembly, CPU Time, Instructions Retired, Wait Time, and Inactive Time. The data shows a range of instructions from 0x6c5600 to 0x6c566d, with CPU times ranging from 0.060s to 3.186s and instructions retired ranging from 67,073 to 12,728,634.

Function Range / Basic Block / Address	Source Line	Assembly	CPU Time	Instructions Retired	Wait Time	Inactive Time
0x6c5600	0	Function range 0x6c5600-0x6c5978	268.288s	1,056,494,654,312		
0x6c5600	0	Block 1	268.157s	1,055,962,887,188		
0x6c5600		mulpd xmm3, xmm0	3.112s	12,427,891,662		
0x6c5604		addpd xmm7, xmm3	0.529s	2,032,037,066		
0x6c5608		movaps xmm3, xmmword ptr [rbx-0x70]	1.812s	7,225,304,386		
0x6c560c		movaps xmm5, xmm3	0.083s	355,717,146		
0x6c560f		mulpd xmm3, xmm0	2.211s	8,971,055,164		
0x6c5613		addpd xmm8, xmm3	1.906s	7,271,515,015		
0x6c5618		movaps xmm3, xmmword ptr [rbx-0x60]	4.500s	18,092,988,610		
0x6c561c		movaps xmm6, xmm3	0.017s	70,273,803		
0x6c561f		mulpd xmm3, xmm0	0.525s	2,054,360,984		
0x6c5623		addpd xmm9, xmm3	0.131s	578,772,055		
0x6c5628		movaps xmm3, xmmword ptr [rbx-0x50]	3.319s	12,921,010,025		
0x6c562c		pshufd xmm1, xmm0, 0x4e	0.016s	84,829,269		
0x6c5631		mulpd xmm0, xmm3	0.379s	1,502,203,392		
0x6c5635		addpd xmm10, xmm0	0.336s	1,450,989,780		
0x6c563a		movaps xmm0, xmmword ptr [rcx+0x10]	3.372s	13,313,573,902		
0x6c563e		mulpd xmm3, xmm1	0.014s	67,073,505		
0x6c5642		addpd xmm14, xmm3	0.362s	1,487,459,978		
0x6c5647		movaps xmm2, xmmword ptr [rcx+0x40]	1.549s	6,173,171,420		
0x6c564b		mulpd xmm4, xmm1	1.442s	5,647,343,685		
0x6c564f		addpd xmm11, xmm4	0.008s	26,962,535		
0x6c5654		movaps xmm3, xmmword ptr [rbx-0x40]	1.602s	6,299,546,943		
0x6c5658		mulpd xmm5, xmm1	0.771s	2,687,860,188		
0x6c565c		addpd xmm12, xmm5	1.024s	3,919,277,667		
0x6c5661		movaps xmm4, xmm3	2.346s	9,619,384,875		
0x6c5664		mulpd xmm6, xmm1	0.368s	1,306,556,762		
0x6c5668		addpd xmm13, xmm6	0.312s	1,350,292,673		
0x6c566d		mulpd xmm3, xmm0	3.186s	12,728,634,324		

Figure 12: Profiling of HPL Start



The screenshot shows the 'Advanced Hotspots' tool interface with the 'Hotspots viewpoint (change)' selected. The 'Assembly' tab is active, displaying a table of assembly instructions grouped by function range. The table includes columns for Function Range / Basic Block / Address, Source Line, Assembly, CPU Time, Instructions Retired, Wait Time, and Inactive Time. The data shows a range of instructions from 0x6c58e9 to 0x6c5968, with CPU times ranging from 0.022s to 4.471s and instructions retired ranging from 99,678 to 12,412,508.

Function Range / Basic Block / Address	Source Line	Assembly	CPU Time	Instructions Retired	Wait Time	Inactive Time
0x6c58e9		addpd xmm13, xmm6	0.060s	208,172,733		
0x6c58ee		mulpd xmm3, xmm2	2.532s	10,430,446,279		
0x6c58f2		addpd xmm7, xmm3	1.309s	5,020,329,310		
0x6c58f6		movaps xmm3, xmmword ptr [rbx+0x50]	2.872s	11,422,369,867		
0x6c58fa		movaps xmm5, xmm3	0.030s	127,392,972		
0x6c58fd		mulpd xmm3, xmm2	0.474s	1,925,583,006		
0x6c5901		addpd xmm8, xmm3	1.810s	6,740,837,318		
0x6c5906		movaps xmm3, xmmword ptr [rbx+0x60]	4.471s	17,750,145,239		
0x6c590a		movaps xmm6, xmm3	0.015s	49,799,094		
0x6c590d		mulpd xmm3, xmm2	0.079s	352,569,637		
0x6c5911		addpd xmm9, xmm3	0.622s	2,373,399,107		
0x6c5916		movaps xmm3, xmmword ptr [rbx+0x70]	3.240s	12,412,508,272		
0x6c591a		pshufd xmm1, xmm2, 0x4e	0.013s	61,982,865		
0x6c591f		mulpd xmm2, xmm3	0.041s	163,087,494		
0x6c5923		addpd xmm10, xmm2	0.866s	3,154,522,922		
0x6c5928		mulpd xmm3, xmm1	3.477s	13,628,252,359		
0x6c592c		addpd xmm14, xmm3	0.013s	82,652,653		
0x6c5931		mulpd xmm4, xmm1	1.353s	5,530,371,836		
0x6c5935		addpd xmm11, xmm4	0.225s	944,011,748		
0x6c593a		add rcx, 0x80	2.291s	9,579,230,090		
0x6c5941		mulpd xmm5, xmm1	0.014s	80,621,660		
0x6c5945		addpd xmm12, xmm5	0.609s	2,289,160,775		
0x6c594a		movaps xmm3, xmmword ptr [rbx+0x80]	2.964s	11,692,748,845		
0x6c5951		mulpd xmm6, xmm1	0.924s	3,700,994,790		
0x6c5955		add rbx, 0x100	0.074s	304,700,689		
0x6c595c		movaps xmm4, xmm3	0.295s	1,109,981,920		
0x6c595f		addpd xmm13, xmm6	2.243s	9,067,366,213		
0x6c5964		sub rax, 0x8	1.318s	5,176,260,723		
0x6c5968		inl 0x6c5600 <Block 1>	0.022s	99,678,836		

Figure 13: Profiling of HPL End

However, we failed to optimize this segment of code because it seems like to have been optimized well. We tended to optimize it by updating the instruction set to SSE4.2

but failed to find a instruction which works. Nor did we find the source code of intel MKL so we can recompile the library.

Finally determined that the computing performance is positively correlated with the value of N (before $N \leq 102400$), and the theoretical optimal value is 135093. The actual optimal value is: 200000. The calculation performance is positively correlated with the value of NB (before $NB \leq 160$), and the performance at $NB = 80$ and $NB = 160$ is equivalent, and the resulting optimal value is 120. $P \times Q = 4 \times 4$ is the value when the performance is optimal. The optimal performance is that it reaches the theoretical performance limit with 84.35%. On our test platform, the optimal parameters are: $N = 20000$, $NB = 120$, $P = 4$, $Q = 4$. The best performance is $71.8Gflops$.

5 HPCG Test

5.1 Hardware Platform

- CPU : Intel Xeon E5 2660v4 (2.0GHz, 14C28T) $\times 2$
- Memory : 32GB DDR4-ECC 2400 MHz $\times 2$
- GPU : Tesla K80 $\times 4$
- Disk : HDD Driver 1TB

5.2 Software Platform

- System : Ubuntu 16.04.3 LTS
- HPCG : HPCG 3.1 binary compiled against CUDA 8.0.61
- MPI : Open MPI 1.10.2
- Compiler : GCC 4.8.5

5.3 Summary

HPCG is designed as a different measure of HPC from HPL.

We setup the environment for HPCG firstly. To compile OpenMPI with CUDA, we configured the source code of openmpi-1.10.2 with option "`-with-cuda=/usr/local/cuda-8.0`".

By comparing through results of different parameters(show in table 5), we found that the performing is best when CPU cores = GPU cores.

Since this is a binary release, optimization on the binary itself is considerably difficult. So we didn't optimize that.

Table 5: HPCG Test Log

Process Grid	Local Domain	Final Performance (GF)
$1 \times 1 \times 1$	$128 \times 128 \times 128$	26.7
$2 \times 1 \times 1$	$128 \times 128 \times 128$	56.2
$2 \times 2 \times 1$	$256 \times 256 \times 256$	109.3
$2 \times 2 \times 1$	$128 \times 128 \times 128$	108
$2 \times 2 \times 2$	$128 \times 128 \times 128$	103.4
$4 \times 2 \times 2$	$128 \times 128 \times 128$	101.2

6 The RELION Test

6.1 Hardware Platform

- CPU : Intel Xeon E5 2660v4 (2.0GHz, 14C28T) $\times 2$
- Memory : 32GB DDR4-ECC 2400 MHz $\times 2$
- GPU : Tesla K80 $\times 4$
- Disk : HDD Driver 1TB

6.2 Software Platform

- Ubuntu 16.04.3 LTS
- CUDA 8.0.61
- Open MPI 1.10.2
- Relion 2.1 Stable

6.3 Parameter Tuning

In this question, we determine that the parameter adjustment is mainly based on the estimating initial noise spectra of STEP1 (Class2D), estimating accuracies in the orthotropy assignment and iteration1 (a total of 25 iterations), and the operating speed, CPU efficiency, and GPU efficiency.

Firstly, when assigning the GPU to the MPI process, the default settings are used. The number of MPI processes must match $4n+1$, 1 is the master process that is responsible for sending and receiving, and the rest is the slave process (responsible for operations).

Secondly, we tried to adjust the number of MPI processes, the total number of threads. But through actual testing found that the effect of this adjustment is not significant. Therefore choose to use 25 MPI processes, each process corresponds to a thread, each card is responsible for 6 MPI process.

Thirdly, trying to adjust the compiler parameters. When using the default cmake (intel compiler) compilation parameters, if you use 25 processes, half of the CPU usage is too low (less than 50%), the GPU also Only reach 80 90W power. If you use 5 processes, CPU usage can be greatly improved, but the calculation speed is extremely low. After adding the O2 compiler parameters, the speed of the three steps can be increased by about one time. At the same time, the O3 compiler parameters are analyzed, and the actual results are not significant.

Finally, the intel compiler includes the parallel parameter by default, and can automatically optimize some loops that can be parallelized. After adding, the running time of estimating initial noise spectra is shortened from 8 ~ 9 seconds to 5 seconds, and there is a significant improvement.

The final optimal operating parameters for STEP1 are as follows:

```
1 mpirun -n 25 "which relion_refine_mpi" --gpu --o Class2D/job007/ --i particles
  .star --ctf --iter 25 --tau2_fudge 2 --particle_diameter 150 --K 100 --
  flatten_solvent --zero_mask --strict_highres_exp 8 --oversampling 1 --
  psi_step 12 --offset_range 5 --offset_step 2 --norm --scale
```

6.4 Code Optimization

6.4.1 First Hot Point

Though we could hardly find a better way like -O2, we have found something ably to better by checking the codes. When we use -O2 to faster the program on Estimating initial noise spectra, we know it could be hardly better on CPU.

So we try to use OpenACC to make the for loop paralleled better running on GPU. And we find a great number of for loops in the codes: fftw.h, fftw.cpp, mask.cpp, matrix1d.h, matrix2d.h.

However, after adding the statements accompanied with their clause:

```
1 #pragma acc kernels
  #pragma acc parallel
```

we find that we could hardly configure the corresponding environment for OpenACC (We have tried GCC5.3.0, GCC6.1.0 and GCC7.1.0). So we're sorry for no results of it.

6.4.2 Second Hot Point

Another point is that, the marcos in multidim_array.h and multidim_array.cpp contains a lot of for loops. But we can't directly add the OpenACC codes in it or add it in front of the quote of the marcos (because the format of the for loop is illegal for OpenACC's optimizing). Our scheme is to rewrite the marcos, move the variables out or to a right place.

6.4.3 Third Hot Point

The 3rd point we have considered to better is that: the optimizing of MPI is based on the parallel of tasks rather than the parallel of calculations. That means, there are a lot of matrixs and FFT operations working on a serial state. At the beginning, we try to use OpenMP to greater the threads and better the efficiency. However, because of our poor hardware, the optimizing isn't evident (actually, we have make all the sources run for the codes). So the next step is to rewrite the logic of MPI's scheduling. However, then we are struggling with the build for OpenACC's environment and decline the time for this way. As a result, we didn't have enough time to achieve it.

6.5 Answer Summary

In this question, the first step and the second step are based on the combination of CPU + GPU computing method, the third step due to the computing platform, using only the CPU to calculate.

The Baidu SkyDrive links to : https://pan.baidu.com/s/1beTeVWpHKK1uIs_qAUywlg. The key is : *nicp*.

The MD5 Code for answer's file is listing as flowing:

- **Step1.Class2D.tar.gz** : 7b64b0357f0dfec9a867438dfd3a605f
- **Step2.Class3D.tar.gz** : 305f41069e20232887ff2f9d8b724261
- **Step3.Refine3D.tar.gz** : 191b5ef232492a74ab346e411e2db68e

6.5.1 Step I : Class 2D

Open in relion_maingui, import ".mrcs" file, node type is 2D micrograph movies (*.mrcs), press run now, press Display to open view. Get figure 14.





Figure 14: Out Put in Step 1

6.5.2 Step II : Class 3D

Open in relion_maingui, import ".mrc" file, node type is 3D reference (*.mrc), press run now, press Display to open the view. We select some pictures from the slice of the stereo image (show in figure 15).


Figure 15: Part of Out Put in Step 2



6.5.3 Step III : Refine 3D

Use UCSF Chimera to adjust. Open the ".mrc" file, then set the step to 1 (connect the initial fractured graph), select level as 0.0272, Radius as 3.3, make the image as close as possible to the original, and then adjust the angle to get the desired image (show in figure 16).

Figure 16: Out Put in Step 3



7 Answer Prediction for Search Query

7.1 Hardware Platform

- CPU : Intel Xeon E5 2660v4 (2.0GHz, 14C28T) \times 2
- Memory : 32GB DDR4-ECC 2400 MHz \times 2
- GPU : Tesla K80 \times 4
- Disk : HDD Driver 1TB

7.2 Software Platform

- Ubuntu 16.04.3 LTS
- CUDA 8.0.61
- CNTK 2.4

7.3 Summary

We learned the usage of CNTK through the baseline code and its documentation, and designed a network with a variety of attention mechanisms. However, due to the constraints of our hardware, the training of the network with self-attention will be very slow so that we cannot adjust the effect of the network. We tried to run test on the MSMARCO dataset after training on the Stanford SQUAD dataset, but the results were worse than baseline. Finally, we used Bi-directional attention flow network architecture then adjust a variety of hyper-parameters and selected the best performing model on the validation set as our submission.

Part of the output is shown below (json format):

```

{"query_id": 0, "answers": ["not"]}
{"query_id": 1, "answers": ["petroleum oil , natural gas , and coal . Fossil
2  fuels , like coal , oil , and natural gas , provide the energy that powers
   our lifestyles and our economy ."]}
{"query_id": 2, "answers": ["The apothem of a regular polygon is a segment
   drawn the center of the polygon ( that is , the center of the circle which
   circumscribes the polygon ) to one side , such that it is perpendicular
   to the side ."]}
{"query_id": 3, "answers": ["$ 600 to $ 1400"]}
4 {"query_id": 4, "answers": ["Hardware is a comprehensive term for all of the
   physical parts of a computer , as distinguished from the data it contains
   or operates on , and the software that provides instructions for the
   hardware to accomplish tasks ."]}
6 {"query_id": 5, "answers": ["a firm that provides service to its customers of
   outsourced ( or Third Party ) logistics services for part , or all of
   their supply chain management functions ."]}
{"query_id": 6, "answers": ["Contrary to the beliefs of those who advocate the
   legalization of marijuana , the current balanced , restrictive , and
   bipartisan drug policies of the United States are working reasonably well
   and they have contributed to reductions in the rate of marijuana use in
   our nation ."]}
8 {"query_id": 7, "answers": ["spiders , scorpions , ticks , mites , harvestmen
   , and their cousins ."]}
{"query_id": 8, "answers": ["southeastern Russia and the Jilin Province of
   northeast China ."]}
10 {"query_id": 9, "answers": ["$ 50 - 75"]}
{"query_id": 10, "answers": ["Proteus mirabilis is part of the normal flora of
   the gastrointestinal tract , and as a result the bacteria enters the
   urinary tract or infects medical equipment by the fecal route ."]}

```

See the appendix for complete results.

References

- [1] <https://software.intel.com/en-us/articles/performance-tools-for-software-developers-hpl-application-note>
- [2] <https://bytefreaks.net/programming-2/>
- [3] <http://www.netlib.org/benchmark/hpl/faqs.html#blsize>
- [4] <https://ark.intel.com>

Appendices

- HPL Log: (path: ./HPL)
- HPCG Log: (path: ./HPCG)
- Answer Prediction For Search Query: (path: ./AnswerPrediction)